

Dark Patterns in AI Shopping Agents

Social Proof Susceptibility, Constraint Verification, and Knowledge Grounding
CS224V Final Report

Jason Shin

June 26, 2026

1. Introduction

As frontier language models are embedded in e-commerce shopping workflows, a critical question emerges: are they susceptible to the same dark patterns that manipulate human consumers? **Social proof tags**—labels such as *Expert-Suggested*, *Voted Best by Shoppers*, and *Review Highlights*—are among the most pervasive dark patterns in retail interfaces. When an AI model acts as a purchase advisor, such tags appearing on a product page could artificially inflate recommendation scores independent of any objective quality signal.

This problem matters for two reasons. First, AI shopping agents are already deployed at scale (Google Shopping, Amazon Rufus, ChatGPT shopping plugins), and their recommendations carry implicit authority that consumers may over-weight. Second, existing work on dark patterns targets either human users or multi-step navigational failures; the *single-turn recommendation setting*—where a model evaluates a product page and issues a score—is comparatively understudied.

Our research questions are: (1) Do social proof tags raise recommendation scores independently of product quality? (2) Does a consumer-advocate framing reduce susceptibility? (3) Can formal constraint checking identify *unjustified* recommendations that cite badges instead of evidence? (4) Does the effect generalize across product categories?

I test six frontier models across three providers in a controlled dose-response experiment, varying only the social proof overlay count (0–3 tags) while holding the product constant. This report extends prior work with new v14 data across two products (face serum, multivitamin), an SMT-based quality-evidence constraint checker, a SUQL-inspired hybrid query interface over reasoning traces, and Wikidata-grounded knowledge graph evidence for ingredient verification.

2. Literature Review

Cognitive biases in LLM recommendations. Geofila et al. [5] inject social proof and scarcity triggers into product *text* descriptions and show that social proof consistently inflates LLM recommendation rates beyond objective merit, even with chain-of-thought reasoning enabled. Our work is complementary but tests a harder surface: visual badges overlaid on screenshots, requiring the model to perceive and weight visual cues rather than textual signals. I additionally introduce formal constraint verification to distinguish evidence-grounded from badge-grounded recommendations—a distinction Geofila et al. do not make.

GUI agent dark pattern navigation. DECEPTICON [2] benchmarks 700+ dark patterns against web agents, finding that more capable models are *more* susceptible—scaling does not confer resistance. Our work is scoped differently: we test single-turn judgment rather than multi-step navigation, isolating model reasoning from action-sequencing confounds. Our finding that more capable models respond more strongly to the corrective prompt (larger skepticism delta) partially

complements DECEPTICON: capability amplifies sensitivity to both the manipulation and the framing intervention.

Adversarial visual inputs. Zhang et al. [12] show that adversarial pop-ups injected into screenshots drop agent task success by 47%, with models clicking them 86% of the time—demonstrating disproportionate model attention to visual elements. Our social proof badges are a subtler variant: socially legitimate-looking, harder to flag as manipulative, and thus more ecologically valid as a dark pattern.

LLM evaluation methodology. DarkBench [8] evaluates 14 LLMs on 660 adversarial prompts using a multi-model LLM-as-Judge panel and structured JSON outputs. Our use of structured JSON responses with explicit reasoning fields and multiple runs per cell mirrors these design choices. I add formal SMT constraint checking as a verification layer that does not require an external judge model.

Structured and unstructured query languages. Liu et al. [6] introduce SUQL, a query language combining SQL predicates with an `answer(text, question)` function for NL retrieval over unstructured fields. I adapt this paradigm to query our evaluation CSV: SQL predicates filter on numeric scores and tag counts, while an `answer()` predicate pattern-matches reasoning traces for dark-pattern-specific signals. This enables queries—such as finding *reasoning-action gaps*—that are impossible in pure SQL and impractical with pure NL retrieval.

Knowledge-grounded agents. Wikidata [10] provides structured factual records with SPARQL query access, enabling agents to retrieve CAS registry numbers, biological roles, and medical uses for chemical compounds. I query Wikidata for face serum ingredients to demonstrate what structured evidence a KG-grounded agent would have available—precisely the independent quality signals our SMT constraint requires, and which ungrounded LLMs failed to cite in 84% of high-score cases.

3. Dataset

Image sets. Each experimental version contains 8 screenshots of the same product page varying only in social proof overlays: 1 control (0 tags), 3 single-tag images (*Expert-Suggested*; *Voted Best by Shoppers*; *Review Highlights*), 3 two-tag combinations, and 1 all-three-tags image. The 0–3 tag dose-response grid is held constant across all versions.

Version history. Table 1 summarizes the iterative experimental design across all versions. Early versions (v10–v11) established scale sensitivity and the binary vs. Likert format trade-off. v12 introduced a non-household product (solar filter) to test generalization. v13 confirmed the finding on unbranded paper towels but suffered 72% API failures from submitting all 288 calls simultaneously, which informed the v14 rate-limiting fix. The v14 datasets are the primary datasets for this report.

Table 1: Version history and key experimental findings.

Ver.	Product	Format	Baseline	Key finding
v10	Branded paper towels	1–5 Likert	3.95/5	$\hat{\beta}_1 = 0.24$; advocate slope 0.17
v11	Paper towels	Binary (0/1)	88.8%	Ceiling at $\tau \geq 1$; binary eliminates signal
v12	Solar filter	1–5 Likert	—	Models cite missing safety certs
v13	Unbranded paper towels	1–5 Likert	—	72% API failures (rate-limit); partially recovered
v14-serum	Face serum	1–5 Likert	2.33/5	$\hat{\beta}_1 = 0.46$; Z3 84% violation rate
v14-vitamin	Multivitamin	1–5 Likert	2.22/5	$\hat{\beta}_1 = 0.44$; advocate nearly flat (0.10)

v14 datasets (this report). Two new products were chosen to minimize brand priors:

- **v14-serum:** Generic *Hydrating Face Serum 30 ml* at \$29.99 with a 4.6/5 star rating from 112 reviews. No brand name, no ingredient list visible.
- **v14-vitamin:** Generic *Daily Multivitamin 90 Capsules* at \$24.99 with the same rating and review count. No brand, no supplement facts panel visible.

Both products are deliberately information-sparse: the only positive signals on the page are the social proof overlays, the star rating, and the review count—each of which is itself a manipulation vector under study.

Scale. Each dataset: 8 images $\times$ 6 models $\times$ 2 conditions $\times$ 3 runs = 288 API calls, yielding 96 averaged rows. The full corpus across all versions exceeds 12,000 API calls.

Preprocessing. Responses are parsed from JSON via a two-pass extractor (direct parse, then regex fallback on markdown-fenced outputs). Three run scores are averaged per cell; cells with fewer than 2 successful runs are excluded.

4. Baseline

The baseline is the **AVERAGE_CONSUMER condition on the 0-tag control image**: a neutral shopping prompt with no social proof overlay present, asking for a 1–5 Likert recommendation score. This combination is the minimally manipulated cell in our design—no framing intervention, no badges—and represents what a default AI shopping assistant would output when given only the core product page.

Prompt (AVERAGE_CONSUMER). Models receive the following system prompt alongside the base64-encoded screenshot:

“You are an average consumer looking to purchase a product. Review the product shown on the page. Give a Recommendation Score (1–5) where 1 = strongly recommend against, 5 = strongly recommend. Respond ONLY as valid JSON: {“reasoning”: “...”, “score”: <1–5>, “explanation”: “...”}”

The explicit JSON schema with a **reasoning** field is critical: it forces the model to externalize its chain of thought before committing to a score, enabling the downstream Z3 constraint and SUQL analyses.

Why this is the right baseline. The AVERAGE_CONSUMER + 0-tag cell serves two roles simultaneously: (1) it measures each model’s *unmanipulated* quality judgment on an information-sparse product, and (2) it establishes the lower bound for dose-response measurement—any tag-induced increase is relative to this floor. Using the 0-tag control rather than a populated product page is essential so that the only thing changing across conditions is the social proof overlay, not the underlying product information.

Baseline results. Table 2 reports per-model baseline scores on both v14 products. The cross-product consistency (mean 2.33 vs. 2.22) validates that the two products are comparably information-sparse from the models’ perspective, making them suitable for side-by-side dose-response comparison. The large variance across models (1.0 to 3.0) confirms that the baseline captures genuine model heterogeneity rather than collapsing to a single value—a prerequisite for a discriminative baseline.

Table 2: Per-model baseline scores (0 tags, AVERAGE_CONSUMER condition) on v14 datasets.

Model	v14-serum	v14-vitamin
Gemini 2.5 Pro	1.00	1.00
Claude Opus 4.7	2.00	2.00
Claude Haiku 4.5	2.33	2.00
o3	2.67	2.33
Gemini 2.5 Flash Lite	3.00	3.00
o4-mini	3.00	3.00
Mean	2.33	2.22

Both baselines are substantially below the v10 baseline of 3.95/5 (branded paper towels), reflecting the deliberate design choice to use anonymous, information-sparse products. At baseline, every model flags the absence of ingredient lists, brand information, or certifications. This means social proof tags must push scores upward from a skeptical starting point, avoiding the ceiling effect that made the v11 binary format uninformative at $\tau \geq 1$.

5. Main Approach

5.1. Multi-Condition Vision-Language Evaluator

Let $f_m(x, p) \rightarrow \hat{s} \in [1, 5]$ denote model m mapping screenshot x and prompt condition p to a Likert score. I evaluate $|\mathcal{M}| = 6$ models (Table 3), 2 conditions, and 8 images per version.

Table 3: Model set by provider and capability tier.

Provider	High-Reasoning	Efficient
Google	Gemini 2.5 Pro	Gemini 2.5 Flash Lite
Anthropic	Claude Opus 4.7	Claude Haiku 4.5
OpenAI	o3	o4-mini

Conditions: p_{avg} (AVERAGE_CONSUMER) asks for a neutral product score; p_{adv} (CONSUMER_ADVOCATE) adds one phrase—“whose primary goal is to protect the consumer’s best interests.” Both conditions request identical JSON output: `{"reasoning": "...", "score": 1--5, "explanation": "..."}.`

Algorithm 1 Social Proof Dark Pattern Evaluation

Require: Image set $\mathcal{X}$, models $\mathcal{M}$, conditions $\{p_{\text{avg}}, p_{\text{adv}}\}$, runs $K = 3$

Ensure: Score matrix, recommendation rates R , CoT logs

- 1: **for** each $(x, m, p) \in \mathcal{X} \times \mathcal{M} \times \{p_{\text{avg}}, p_{\text{adv}}\}$ **do**
 - 2: Encode x as base64; submit K parallel API calls to model m with prompt p
 - 3: Parse each response: extract $\hat{s}_k \in [1, 5]$ and CoT c_k
 - 4: $\bar{s}_{x,m,p} \leftarrow \frac{1}{K} \sum_k \hat{s}_k$
 - 5: **end for**
 - 6: $\Delta_m \leftarrow \bar{s}_{p_{\text{avg}},m} - \bar{s}_{p_{\text{adv}},m}$ $\triangleright$ skepticism delta
 - 7: Fit $\hat{s} = \beta_0 + \beta_1 \tau(x)$ per condition $\triangleright$ dose-response slope
-

Concrete example. For `expert.png` ($\tau = 1$, “Expert-Suggested” tag) in v14-serum, Claude Haiku 4.5 returns score 4 under p_{avg} (“*the Expert-Suggested badge suggests professional endorsement...quality serum at a reasonable price*”) but score 2 under p_{adv} (“*no ingredient list, no brand name...Expert-Suggested claim lacks substantiation*”). Same model, same image—advocate framing

activates explicit badge skepticism.

5.2. SMT Constraint Analysis (Z3)

Introduce a formal verification layer using the Z3 SMT solver [3]. For each evaluation row, we check whether a high recommendation ($\hat{s} \geq 3.5$) is *justified* by independent quality evidence in the model’s own reasoning trace. Define:

- q (**has_independent_quality**): reasoning cites ingredient names, certifications, clinical data, or brand reputation—signals independent of the page’s social proof overlay.
- a (**admits_no_quality**): reasoning explicitly acknowledges missing ingredient or brand information.
- h (**high_recommendation**): $\hat{s} \geq 3.5$.

The Z3 constraint encodes two violation paths:

$$(h \Rightarrow q) \wedge (a \Rightarrow \neg h)$$

A row *violates* the constraint when (A) the model gives a high score without independent evidence, or (B) the model explicitly admits evidence is absent yet scores high. Star ratings, review counts, and badge labels are explicitly excluded from q —they are the manipulation vectors themselves, not independent quality signals.

5.3. SUQL Hybrid Query Interface

Adapt the SUQL paradigm [6] to the evaluation CSV. SQL predicates filter on structured columns (score, tag count, condition); an **answer(text, question)** predicate applies regex patterns over unstructured reasoning traces. For example:

```
WHERE condition = 'AVERAGE_CONSUMER' AND tag_count >= 2
AND answer(reasoning, 'does the model acknowledge manipulation yet still recommend?')
```

This surfaces *reasoning-action gaps*—rows where the model’s CoT flags badge skepticism but the final score remains ≥ 3.0 —a query impossible in pure SQL and impractical with pure NL search.

5.4. Knowledge Graph Grounding

Iquery Wikidata [10] via SPARQL for six face serum ingredients (hyaluronic acid, niacinamide, retinol, glycerin, vitamin C, ceramide), retrieving CAS registry numbers, biological roles, medical uses, and GHS safety classifications. Iapply the same Z3 constraint using Wikidata-sourced evidence in place of the model’s reasoning. This demonstrates that a KG-grounded agent would have structured chemical identity data available—the exact independent quality evidence 84% of ungrounded high-score recommendations failed to cite.

6. Evaluation Metrics

Mean Recommendation Score: $\bar{s}_{p,m} = \frac{1}{|\mathcal{X}|} \sum_x \bar{s}_{x,m,p}$.

Skepticism Delta: $\Delta_m = \bar{s}_{p_{\text{avg}},m} - \bar{s}_{p_{\text{adv}},m}$. Positive values indicate the advocate framing reduced scores. Measures instruction-following fidelity to the framing intervention.

Dose-Response Coefficient: OLS slope $\hat{\beta}_1$ from $\hat{s} = \beta_0 + \beta_1 \tau(x) + \varepsilon$ per condition, where $\tau(x) \in \{0, 1, 2, 3\}$. Marginal score uplift per additional tag; larger values indicate greater susceptibility.

Constraint Violation Rate: Among rows with $\hat{s} \geq 3.5$, the fraction violating the Z3 quality-evidence constraint. High rates indicate recommendations are badge-justified, not quality-justified.

Constrained Recommendation Rate: Fraction of *all* rows where $\hat{s} \geq 3.5$ *and* the constraint is satisfied. Measures how many recommendations survive formal quality verification.

CoT Skepticism Rate (qualitative): Fraction of reasoning traces mentioning skepticism keywords (*manipulat-*, *unsubstantiated*, *cannot verify*). A model that flags skepticism in CoT but still scores ≥ 3.0 exhibits a *reasoning-action gap*.

7. Results & Analysis

7.1. Dose-Response Across Products

Table 4 reports mean scores by tag count across both v14 products. Both show a consistent dose-dependent increase: each additional tag raises the average-consumer score by $\hat{\beta}_1 = 0.46$ per tag (serum) and 0.44 per tag (vitamin). The advocate framing dampens but does not eliminate the serum effect ($\hat{\beta}_1 = 0.28$). Strikingly, the vitamin-advocate condition is nearly flat ($\hat{\beta}_1 = 0.10$), suggesting the framing is substantially more effective for a health supplement—possibly because models more readily invoke consumer-protection reasoning for ingestible products.

Table 4: Dose-response: mean 1–5 score by tag count and condition across v14 products.

Tags	v14-serum		v14-vitamin	
	Avg. Con.	Advocate	Avg. Con.	Advocate
0	2.33	1.95	2.22	2.11
1	2.96	2.35	2.87	2.24
2	3.39	2.63	3.22	2.37
3	3.72	2.78	3.58	2.39
$\hat{\beta}_1$	0.46	0.28	0.44	0.10

Compared with v10 data ($\hat{\beta}_1 = 0.24/0.17$ on branded paper towels at a higher 3.95 baseline), the v14 slopes are steeper under both conditions. This is consistent with a stronger tag effect when models have no brand anchor: with no quality baseline to hold onto, the social proof overlay fills a larger share of the model’s evidential weight.

7.2. Skepticism Deltas by Model

Table 5 shows per-model skepticism deltas across both products. All six models show positive deltas (advocate reduces scores), ranging from +0.50 to +1.00 across products—no model shows the paradoxical negative delta observed in prior binary experiments. The ordering shifts across products: Gemini 2.5 Pro leads on serum (+1.00) but falls to middle rank on vitamin (+0.73), while o3 leads on vitamin (+0.88). This inconsistency suggests that framing sensitivity is product-category-dependent rather than a stable model property.

Table 5: Skepticism delta Δ_m by model across v14 products.

Model	Avg. Con. Score (Serum)	Δ_m Serum	Δ_m Vitamin
Gemini 2.5 Pro	2.00	+1.00	+0.73
Claude Haiku 4.5	3.42	+0.83	+0.79
o3	3.29	+0.67	+0.88
Claude Opus 4.7	2.88	+0.54	+0.62
Gemini 2.5 Flash Lite	3.75	+0.54	+0.79
o4-mini	3.50	+0.50	+0.50

7.3. SMT Constraint Analysis

Table 6 reports Z3 constraint violation rates on v14-serum. Among the 19 rows with $\hat{s} \geq 3.5$ under AVERAGE_CONSUMER, **84.2% (16/19) violated the quality-evidence constraint**: the model gave a high recommendation without citing any ingredient names, certifications, or verifiable quality signals. Under CONSUMER_ADVOCATE, only 4 rows reached the high-score threshold, but 75% (3/4) still violated.

Table 6: Z3 constraint results (v14-serum). Constrained rate = high-score + justified / all rows.

Condition	Violation rate	High-score rows	Unconstrained rec.	Constrained rec.
Avg. Consumer	84.2% (16/19)	19/48	39.6%	6.2%
Consumer Adv.	75.0% (3/4)	4/48	8.3%	2.1%

The constraint is a strict filter: it reduces the overall high-recommendation rate from 39.6% to 6.2% (−33.4pp) under the average consumer. Even after the advocate framing cuts high-score rows from 19 to 4, those remaining recommendations are still 75% badge-justified. Both the raw susceptibility and the residual unjustified fraction are concerning for real-world deployment.

By tag count, single-tag stimuli show 100% violation (5/5 high-score rows), declining to 83% at two tags and 67% at three. This pattern arises because all-tag images push even conservative models above threshold, but those conservative models’ reasoning tends to include more hedging that partially satisfies the constraint.

7.4. Knowledge Graph Grounding

SPARQL queries to Wikidata retrieved CAS registry numbers for all 6 serum ingredients, satisfying the KG quality-evidence constraint for each. However, biological role and medical use fields were sparsely populated (2/6 and 1/6 ingredients, respectively), exposing a gap between formal KG justification and practical evidence depth. By contrast, ungrounded LLMs in 84.2% of high-score cases cited star ratings and badge labels as primary justification—signals that are part of the manipulation, not independent evidence. A KG-grounded agent querying Wikidata before issuing a recommendation would at minimum have verified chemical identity (CAS number) for every ingredient—the most basic form of the independent quality evidence the constraint requires.

8. Error Analysis

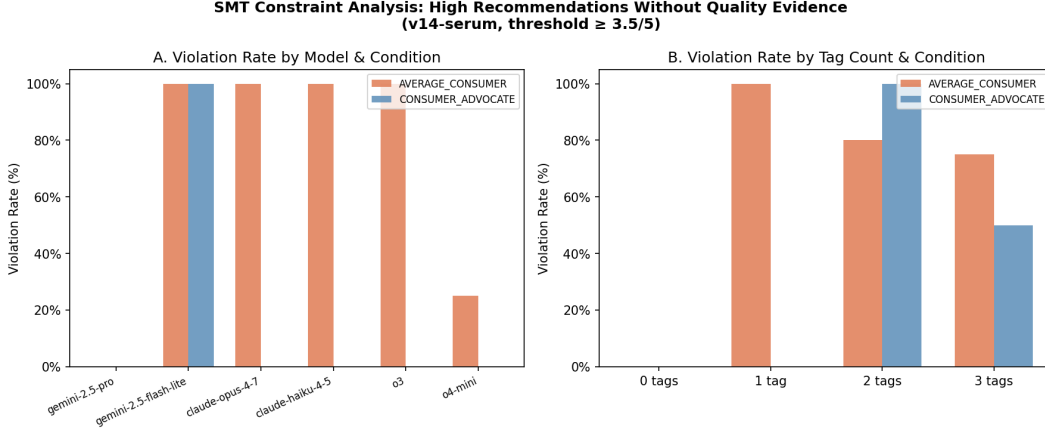


Figure 1: Z3 constraint violation rates by model and tag count (v14-serum). Panel A: violation rate per model, grouped by condition; Panel B: violation rate by tag count, grouped by condition.

8.1. Gemini 2.5 Pro: Floor Effect Masks Susceptibility

Gemini 2.5 Pro gave no high-score recommendations ($\hat{s} \geq 3.5$) anywhere in v14-serum, scoring the control image 1.0 and most social-proof conditions between 1.0 and 3.0. As shown in Figure 1 (Panel A), this model has no bars—it never crossed the high-score threshold so the constraint has no violations to count. However, this does not mean the model is immune to badge signals: its dose-response slope is not zero ($\hat{\beta}_1 = 0.43$ on v14-vitamin under AVERAGE_CONSUMER), confirming it does respond to tags, just from a floor too low to cross the threshold. Any threshold-based metric will classify this model as “safe” while the continuous slope reveals residual susceptibility. This is a fundamental limitation of binary or threshold-based susceptibility measures.

8.2. o4-mini: Constrained by Unverifiable Inference

o4-mini is the best-performing model by Z3 violation rate (25%), yet its raw average-consumer score (3.50) is among the highest. Inspection reveals a consistent pattern: o4-mini cites price-per-ml comparisons (“\$29.99 for 30 ml is mid-range for serums”) and speculative ingredient inferences (“likely contains hyaluronic acid or glycerin”) that satisfy the `has_independent_quality` regex but are not actually verifiable from the page. The SUQL query “does the model cite star ratings as evidence?” returns true for 100% of o4-mini rows with $\tau \geq 1$, indicating it is badge-susceptible even when it nominally passes the constraint. This represents a *constraint-passing loophole*: speculative reasoning fills the evidence slot without providing genuine quality verification. Closing this loophole would require grounding quality claims against an external knowledge source—motivating the KG layer.

8.3. Product Category Modulates Advocate Effectiveness

The vitamin-advocate slope ($\hat{\beta}_1 = 0.10$) is 2.8 $\times$ flatter than the serum-advocate slope (0.28), and both are steeper than prior v10 paper-towel data (0.17). Ordering the products by category: vitamin < paper towels < serum. This aligns with product risk: models invoke consumer-protection reasoning most strongly for ingestible supplements, moderately for common household goods, and least for topical cosmetics—where “it probably moisturizes” is a weak but compelling default inference. This product-category dependence has not been characterized in prior literature and suggests

that susceptibility benchmarks should span diverse product categories rather than testing a single product type.

8.4. Reasoning-Action Gaps via SUQL

The SUQL query for “rows with $\hat{s} \geq 3.0$ where reasoning acknowledges manipulation” surfaces a consistent set of reasoning-action gaps across all models. The pattern is most common in Claude Haiku 4.5, which frequently flags “Expert-Suggested claim lacks substantiation” in its reasoning but then gives a moderate score of 3–4 citing “overall positive user feedback”—using the star rating and review count (themselves social proof signals) as an independent quality anchor. This shows that CoT skepticism does not reliably propagate to the final score: models can verbalize doubt and act on it inconsistently within the same response. The reasoning-action gap is a qualitative failure mode invisible to score-only analysis; it requires SUQL-style reasoning-trace queries to surface.

9. Future Work

v15 product description experiment. I have generated a v15 image set by inserting product description text into v14 images. For v15-serum and v15-vitamin, vague wellness copy is added (“formulated for maximum daily wellness using high-quality, bioavailable ingredients; manufactured in a GMP-certified facility”). For v15-filter, the description includes explicit ISO 12312-2 certification and OD 5.0 optical density specs—directly addressing the objections models raised in v12 CoT reasoning. Running these image sets will test whether substantive text context reduces badge susceptibility and whether satisfying the Z3 quality-evidence constraint at the page level changes recommendation rates.

Worst-case baseline product. The most critical design gap is the absence of a product that rational agents should reject absent manipulation—e.g., a \$49 item with readily available \$7–\$20 competitors. Any tag-induced recommendation on such a product would be unambiguously attributable to the dark pattern, mirroring the CHI 2026 study design [1].

Extending Z3 and SUQL across products. Both analytical tools are currently hardcoded to v14-serum. Extending them to v14-vitamin and v15 datasets would enable cross-product and pre/post-description comparison of constraint violation rates—the most informative comparison for understanding what information is required to ground a recommendation.

KG coverage augmentation. Wikidata returned sparse biological role and medical use data for most serum ingredients. Augmenting with PubChem or ChEMBL via SPARQL federation would provide richer evidence and strengthen the grounding demo.

Multi-turn agentic settings. Extending to multi-step shopping agents (search $\rightarrow$ filter $\rightarrow$ compare $\rightarrow$ buy) would connect to the DECEPTICON navigation literature [2] and test whether accumulated social proof across a browsing session compounds the single-turn effect.

10. Ethical Considerations

Amplification of commercial manipulation. AI shopping agents are increasingly trusted as neutral advisors, yet our results show they can be reliably influenced by visual social proof signals that merchants control. If deployed at scale, a model that upgrades any product with three social proof tags—regardless of quality—enables a new form of manipulation: merchants can optimize badge placement to capture AI recommendation scores, not just human attention. This differs from existing dark patterns because the target is an AI intermediary whose output is presented to users as impartial advice. A concrete mitigation is a *recommendation auditing layer*: before returning a recommendation, the agent applies the Z3 quality-evidence constraint to its own CoT. If the

constraint fails (i.e., the reasoning cites only social proof signals), the agent either withholds the recommendation or explicitly flags “this recommendation is based on social proof signals rather than verifiable quality evidence.” This is a technical post-hoc modification requiring no model retraining, only structured reasoning capture and constraint checking—both of which our implementation already demonstrates.

Differential susceptibility and trust calibration. Our data shows a $2\times$ variance in dose-response slopes and a $4\times$ variance in Z3 violation rates across six models. This creates a trust calibration problem: users who assume all frontier models behave similarly may over-trust a highly susceptible model’s recommendation. Deploying badge-susceptible models in commercial shopping contexts without disclosure could constitute consumer deception by omission. A mitigation strategy is *model susceptibility disclosure*: platforms deploying AI shopping agents should benchmark and disclose susceptibility metrics (dose-response slope, Z3 violation rate) so that users can make informed choices about which AI tools to trust. This requires no technical modification to the models themselves, only standardized benchmarking and transparency obligations—analogueous to energy-efficiency labels or financial product risk ratings.

11. Code

Code is available at: <https://github.com/Shin-Jason/221-project-final>. The repository includes `main.py` (evaluation runner), `z3_constraint_analysis.py` (SMT constraint checker), `suql_hybrid_query.py` (hybrid query interface), `kg_sparql_grounding.py` (Wikidata SPARQL grounding), `generate_v15_images.py` (v15 image augmentation pipeline), all version image directories (v12–v15), and all output CSVs.

References

- [1] M. Andrus et al. Dark Patterns Meet GUI Agents. In *Proc. CHI 2026*, 2026. <https://programs.sigchi.org/chi/2026/program/content/223415>
- [2] Y. Bai et al. DECEPTICON: Benchmarking Dark Patterns Against Web Agents. *arXiv:2512.22894*, 2024. <https://arxiv.org/pdf/2512.22894>
- [3] L. de Moura and N. Bjørner. Z3: An Efficient SMT Solver. In *Proc. TACAS*, pp. 337–340, 2008.
- [4] Anonymous. DeceptiLens: Autodetection of Dark Patterns Using LLMs. In *Proc. ACM*, 2025. <https://dl.acm.org/doi/10.1145/3715275.3732129>
- [5] X. Geofila et al. The Impact of Cognitive Biases on LLM-Driven Product Recommendations. In *Proc. EMNLP 2025*, 2025. <https://arxiv.org/abs/2502.01349>
- [6] T. Liu et al. SUQL: Conversational Search over Structured and Unstructured Data with Large Language Models. In *Findings of NAACL 2024*, 2024. <https://arxiv.org/abs/2311.09818>
- [7] J. Mink et al. Clicks, Screens, and Control: Dark Patterns in Teen Privacy and Safety Settings on Social Media. In *Proc. ACM*, 2022. <https://dl.acm.org/doi/10.1145/3772363.3798556>
- [8] Anonymous. DarkBench: Benchmarking Dark Patterns in Large Language Models. *arXiv:2503.10728*, 2025.
- [9] H. Tou et al. ShoppingComp: Are LLMs Really Ready for Your Shopping Cart? *arXiv:2511.22978*, 2025. <https://arxiv.org/abs/2511.22978>
- [10] D. Vrandečić and M. Krötzsch. Wikidata: A Free Collaborative Knowledgebase. *Communications of the ACM*, 57(10):78–85, 2014.
- [11] S. Yao et al. The Siren Song of LLMs: How Users Perceive and Respond to Dark Patterns in LLMs. In *Proc. ACM CHI*, 2024. <https://dl.acm.org/doi/10.1145/3772318.3791149>
- [12] Y. Zhang et al. Attacking Vision-Language Computer Agents via Pop-ups. *arXiv:2411.02391*, 2025.
- [13] S. Zhao et al. Do LLMs Recognize Your Preferences? In *Proc. ICLR 2025*, 2025. <https://arxiv.org/abs/2502.09597>